

```
#1.Download the Employee Attrition Dataset
#https://www.kaggle.com/datasets/patelprashant/employee-attrition
#2.Perform Data Preprocessing
#3.Model Building using Logistic Regression and Decision Tree
#4.Calculate Performance metrics
```

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

df = pd.read_csv('Employee-Attrition.csv')

# Check for missing values
print(df.isnull().sum())

# Check for outliers
import seaborn as sns
sns.boxplot(x = 'Age', y = 'Attrition', data = df)
sns.boxplot(x = 'DailyRate', y = 'Attrition', data = df)

# Check data types
print(df.dtypes)
```

```

RelationshipSatisfaction    0
StandardHours              0
StockOptionLevel           0
TotalWorkingYears          0
TrainingTimesLastYear      0
WorkLifeBalance            0
YearsAtCompany             0
YearsInCurrentRole         0
YearsSinceLastPromotion    0
YearsWithCurrManager       0
dtype: int64
Age                        int64
Attrition                 object
BusinessTravel            object
DailyRate                int64
Department               object
DistanceFromHome          int64
Education                int64
EducationField            object
EmployeeCount             int64
EmployeeNumber            int64
EnvironmentSatisfaction   int64
Gender                   object
HourlyRate               int64
JobInvolvement            int64
JobLevel                 int64
JobRole                  object
JobSatisfaction           int64
MaritalStatus            object
MonthlyIncome            int64
MonthlyRate              int64
NumCompaniesWorked       int64
Over18                   object
OverTime                 object
PercentSalaryHike        int64
PerformanceRating        int64
RelationshipSatisfaction  int64
StandardHours            int64
StockOptionLevel         int64
TotalWorkingYears        int64
TrainingTimesLastYear    int64
WorkLifeBalance          int64
YearsAtCompany           int64

```

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Handle missing values
df = df.dropna()

# Handle outliers
df = df[(df['Age'] < 70) & (df['DailyRate'] < 200000)]

# Encode categorical features
df['Gender'] = df['Gender'].map({'Male': 0, 'Female': 1})
df['Department'] = df['Department'].map({'Sales': 0, 'Engineering': 1, 'Marketing': 2})

# Scale the numeric features
scaler = StandardScaler()
X = df[['Age', 'DailyRate', 'Gender', 'Department']]
X = scaler.fit_transform(X)

# Split the data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, df['Attrition'], test_size=0.25, random_state=42)

# Convert the NumPy arrays to Pandas DataFrames
X_train_df = pd.DataFrame(X_train)
X_test_df = pd.DataFrame(X_test)

# Save the preprocessed data to CSV files
X_train_df.to_csv('X_train.csv', index=False)
X_test_df.to_csv('X_test.csv', index=False)
y_train.to_csv('y_train.csv', index=False)
y_test.to_csv('y_test.csv', index=False)

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/utils/extmath.py:1047: RuntimeWarning: invalid value encountered in divide
  updated_mean = (last_sum + new_sum) / updated_sample_count
/usr/local/lib/python3.10/dist-packages/sklearn/utils/extmath.py:1052: RuntimeWarning: invalid value encountered in divide
  T = new_sum / new_sample_count

```

```
/usr/local/lib/python3.10/dist-packages/sklearn/utils/extmath.py:1072: RuntimeWarning: invalid value encountered in divide
  new_unnormalized_variance -= correction**2 / new_sample_count
```

```
X_train_df.to_csv
```

```
<bound method NDFrame.to_csv of
0    -0.170973 -1.235942 NaN NaN
1    -0.503476 -0.487829 NaN NaN
2     2.489050  1.055620 NaN NaN
3    -0.946813  1.676977 NaN NaN
4     1.048204  0.807078 NaN NaN
...
329 -0.946813 -1.233456 NaN NaN
330 -0.170973  1.023310 NaN NaN
331 -0.835979 -0.823361 NaN NaN
332 -0.281807 -0.239286 NaN NaN
333 -0.392642 -1.119127 NaN NaN

[334 rows x 4 columns]>
```

```
X_test_df.to_csv
```

```
<bound method NDFrame.to_csv of
0     0.494033 -1.352757 NaN NaN
1     1.269873 -0.229345 NaN NaN
2     0.494033 -0.385926 NaN NaN
3     1.824044  0.913951 NaN NaN
4    -0.835979  1.018339 NaN NaN
...
107 -0.170973 -0.092646 NaN NaN
108  0.272364  0.235430 NaN NaN
109 -0.725144  1.338959 NaN NaN
110  0.937370 -1.228485 NaN NaN
111 -0.946813  0.923893 NaN NaN

[112 rows x 4 columns]>
```

```
y_train.to_csv
```

```
<bound method NDFrame.to_csv of 1281    Yes
935      No
70       No
1369    Yes
433      No
...
374      No
888      No
1172     No
1446     No
363     Yes
Name: Attrition, Length: 334, dtype: object>
```

```
y_test.to_csv
```

```
<bound method NDFrame.to_csv of 951    No
1204    Yes
403     No
1396    Yes
265     No
...
1167    Yes
33     Yes
167     No
504     Yes
1337    No
Name: Attrition, Length: 112, dtype: object>
```

```
#3
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
```

```
# Load the preprocessed training data
X_train = pd.read_csv('X_train.csv')
y_train = pd.read_csv('y_train.csv')['Attrition']
```

```
# Load the preprocessed test data
```

```

X_test = pd.read_csv('X_test.csv')
y_test = pd.read_csv('y_test.csv')['Attrition']

# Create a logistic regression model
logistic_model = LogisticRegression(random_state=42)

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.impute import SimpleImputer

# Impute the missing values in the Age feature
imputer = SimpleImputer(strategy='mean')
df['Age'] = imputer.fit_transform(df[['Age']])

# Split the data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(df[['Age', 'DailyRate', 'Gender', 'Department']], df['Attrition'], test_size=0.25, random

# Create a logistic regression model
logistic_model = LogisticRegression(random_state=42)

# Train the logistic regression model
logistic_model.fit(X_train, y_train)

# Evaluate the logistic regression model on the test data
y_pred = logistic_model.predict(X_test)
accuracy = logistic_model.score(X_test, y_test)
print('Accuracy:', accuracy)

```

```

-----
ValueError                                Traceback (most recent call last)
<ipython-input-23-63d8f7aba945> in <cell line: 17>()
    15
    16 # Train the logistic regression model
--> 17 logistic_model.fit(X_train, y_train)
    18
    19 # Evaluate the logistic regression model on the test data

----- 4 frames -----
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py in _assert_all_finite(X,
allow_nan, msg_dtype, estimator_name, input_name)
    159         "#estimators-that-handle-nan-values"
    160     )
--> 161     raise ValueError(msg_err)
    162
    163

ValueError: Input X contains NaN.
LogisticRegression does not accept missing values encoded as NaN natively. For supervised learning,
you might want to consider sklearn.ensemble.HistGradientBoostingClassifier and Regressor which accept
missing values encoded as NaNs natively. Alternatively, it is possible to preprocess the data, for
instance by using an imputer transformer in a pipeline or drop samples with missing values. See
https://scikit-learn.org/stable/modules/impute.html You can find a list of all estimators that handle
NaN values at the following page: https://scikit-learn.org/stable/modules/impute.html#estimators-that-handle-nan-values

```

```

from sklearn.impute import SimpleImputer

# Impute the missing values in the Age feature of the train set
imputer = SimpleImputer(strategy='mean')
X_train['Age'] = imputer.fit_transform(X_train[['Age']])

# Impute the missing values in the Age feature of the test set
imputer = SimpleImputer(strategy='mean')
X_test['Age'] = imputer.fit_transform(X_test[['Age']])

```